

Adarsh Dwivedi

Agentic AI Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | linkedin.com/in/adarshdwivedi30 | github.com/adarshcod30 | hackerrank.com/profile/23ucs509

Final-year CS engineer who builds agents that are allowed to propose but never to decide alone. At Deloitte, moved a live SAP configuration agent onto a self-hosted agent server and proved it on the real SAP tenant, completing 8 of 8 tasks with two faults recovered mid-run. Ranked #12 in HackerRank Orchestrate after climbing from #427 across three 24-hour challenges, and built a permission-gated MCP tool layer of 19 tools that no agent request escapes.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 - Aug 2026

AI Engineer (Agentic AI) Intern

Gurugram, India

- Migrated SAP Agentic, Deloitte's system that automates SAP configuration, from LangGraph running inside a FastAPI request handler onto a self-hosted Aegra agent server, with threads, runs and checkpoints in PostgreSQL inside the firewall at no licence cost.
- Proved the new path on the live SAP tenant: a Bank Accounting goal was planned into 8 configuration tasks, executed in a real browser through Playwright MCP recovered from a duplicate-entry conflict and a mis-filled field, and finished 8 of 8 in about 15 minutes.
- Replaced request-time graph assembly with a graph factory that registers one static graph and carries per-request variation in run state, verifying the migrated output identical to the original field by field behind an unchanged API.

PROJECTS

AGENTIQ, Agentic API Testing and Security Validation (B.Tech final-year project) | [Live](#) | [GitHub](#)

Node.js, MCP, LLM Agents, Zod, Vitest

- Functional tests live in Postman and security tests in ZAP, and the moment an LLM acts against a live endpoint nobody can account for what it did. Built one platform where agents turn a URL and plain-English intent into executable multi-assertion tests.
- Constrained every agent action to an MCP tool layer of 19 Zod-schema'd tools across 6 risk classes, with per-host grants, a filesystem jail, a process sandbox and an SSRF guard, so no outbound request escapes policy or audit.
- Probed 8 families mapped to the OWASP API Security Top 10 against a benign baseline, so a finding needs a material deviation rather than a suspicious-looking string, and backed the system with 674 tests in CI.

Travel Planner Agent, Five Generations of a Multi-Agent Planner | [GitHub](#)

LangGraph, MCP, Playwright, Bedrock, Aegra

- Built five successive generations of the same planner, linear, parallel, orchestrator, human-in-the-loop and an MCP factory graph, keeping every generation runnable so the architectural trade-offs can be compared rather than argued about.
- Served all five concurrently through Aegra on AWS Bedrock with a standalone MCP server exposing 10 travel tools, and covered the backend with 630 tests against a real Postgres.
- Wired live Playwright browser automation the user can take over mid-run, so the agent hands control back at exactly the point where a human should decide.

Buy or Wait, Financial Decision Agent (HackerRank Orchestrate Sep 2026, ranked #12) | [GitHub](#)

Python, Amazon Bedrock, Agent Loop

- Ranked #12 of a 24-hour agentic challenge by letting the agent propose and a deterministic layer decide: a 90-day cash-flow forecast keeps only payment plans that never breach the minimum balance.
- Re-applied the deterministic ranking after every model call so an LLM choice cannot silently alter a graded field, and verified each rendered row against the submission contract with self-repair, emitting the next safe plan when the top one fails.
- Made the whole run reproducible from cold: unzipping the submission with every AWS variable cleared regenerates the output byte for byte, so the result does not depend on the machine that produced it.

TECHNICAL SKILLS

Agentic AI: LangGraph, MCP (Model Context Protocol), Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, Agent Protocol, Aegra, Playwright MCP, LLMops

Generative AI: LangGraph, LangChain, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Backend and APIs: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Infrastructure: Docker, AWS, Google Cloud Run, Vercel, Render, GitHub Actions, CI/CD, Linux, Git, Observability, Reproducible Builds

Languages: Python, C/C++ , Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- HackerRank Orchestrate, three 24-hour agentic AI challenges: ranked #12 with a silver medal, top 1% (Sep 2026), #290 with a bronze medal, top 15% (Aug 2026), and #427, top 25% (Jun 2026).
- Competition submissions, each public on GitHub: [Buy or Wait, #12](#) · [Message Notification Router, #290 bronze](#) · [Multi-Modal Evidence Review, #427](#) · [Artifact Repro Triage, micro1](#)

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 - 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.