

Adarsh Dwivedi

Agentic AI Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | linkedin.com/in/adarshdwivedi30 | github.com/adarshcod30 | hackerrank.com/profile/23ucs509

Final-year CS engineer who builds agents that are allowed to propose but never to decide alone. Shipped agentic systems at Deloitte, taking LangGraph agents from local scripts to production services with human-in-the-loop approval on a live enterprise SAP workflow. Ranked #12 in HackerRank Orchestrate after climbing from #427 across three 24-hour challenges, and built a permission-gated MCP tool layer of 19 tools that no agent request escapes.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 - Aug 2026

AI Engineer (Agentic AI) Intern

Gurugram, India

- Took LangGraph agents from local scripts to network-reachable production services holding long-term memory in Postgres and surviving restarts, validating the self-hosted deployment architecture before the team committed engineering time to it.
- Root-caused a silent startup failure where an async Postgres driver deadlocked against Windows network I/O scheduling, evaluated every workaround against the team's constraints, and selected the foundation the rest of the migration was built on.
- Authored the migration plan for a production agent driving a live enterprise SAP workflow end to end with human-in-the-loop approval, defined a parity checklist, and validated both execution paths against it before either was replaced.

PROJECTS

AGENTIQ, Agentic API Testing and Security Validation (B.Tech final-year project) | [Live](#) | [GitHub](#)

Node.js, MCP, LLM Agents, Zod, Vitest

- Functional tests live in Postman and security tests in ZAP, and the moment an LLM acts against a live endpoint nobody can account for what it did. Built one platform where agents turn a URL and plain-English intent into executable multi-assertion tests.
- Constrained every agent action to an MCP tool layer of 19 Zod-schema'd tools across 6 risk classes, with per-host grants, a filesystem jail, a process sandbox and an SSRF guard, so no outbound request escapes policy or audit.
- Probed 8 families mapped to the OWASP API Security Top 10 against a benign baseline, so a finding needs a material deviation rather than a suspicious-looking string, and backed the system with 674 tests in CI.

Travel Planner Agent, Five Generations of a Multi-Agent Planner | [GitHub](#)

LangGraph, MCP, Playwright, Bedrock, Aegra

- Built five successive generations of the same planner, linear, parallel, orchestrator, human-in-the-loop and an MCP factory graph, keeping every generation runnable so the architectural trade-offs can be compared rather than argued about.
- Served all five concurrently through Aegra on AWS Bedrock with a standalone MCP server exposing 10 travel tools, and covered the backend with 630 tests against a real Postgres.
- Wired live Playwright browser automation the user can take over mid-run, so the agent hands control back at exactly the point where a human should decide.

Buy or Wait, Financial Decision Agent (HackerRank Orchestrate Sep 2026, ranked #12) | [GitHub](#)

Python, Amazon Bedrock, Agent Loop

- Ranked #12 of a 24-hour agentic challenge by letting the agent propose and a deterministic layer decide: a 90-day cash-flow forecast keeps only payment plans that never breach the minimum balance.
- Re-applied the deterministic ranking after every model call so an LLM choice cannot silently alter a graded field, and verified each rendered row against the submission contract with self-repair, emitting the next safe plan when the top one fails.
- Made the whole run reproducible from cold: unzipping the submission with every AWS variable cleared regenerates the output byte for byte, so the result does not depend on the machine that produced it.

TECHNICAL SKILLS

Agentic AI: LangGraph, MCP (Model Context Protocol), Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, LLMops, DeepEval, Ragas, Langfuse, Arize Phoenix

Generative AI: LangGraph, LangChain, LangSmith, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Backend and APIs: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Infrastructure: Docker, AWS, Google Cloud Run, Vercel, Render, GitHub Actions, CI/CD, Linux, Git, Observability, Reproducible Builds

Languages: Python, C/C++ , Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- HackerRank Orchestrate, three 24-hour agentic AI challenges: ranked #12 with a silver medal, top 1% (Sep 2026), #290 with a bronze medal, top 15% (Aug 2026), and #427, top 25% (Jun 2026).
- Competition submissions, each public on GitHub: [Buy or Wait, #12](#) · [Message Notification Router, #290 bronze](#) · [Multi-Modal Evidence Review, #427](#) · [Artifact Repo Triage, micro1](#)

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 - 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.