

Adarsh Dwivedi

AI Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | [linkedin.com/in/adarshdwivedi30](https://www.linkedin.com/in/adarshdwivedi30) | github.com/adarshcod30

Final-year CS engineer shipping production AI with evaluation wired in from the start. At Deloitte, took LangGraph agents to network-reachable production services and benchmarked four evaluation frameworks against the team's reliability criteria. Across 32 public repositories, most of them deployed and free to use, every reported number is reproducible from the repository, and three of them ship as installable packages on PyPI.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 - Aug 2026

AI Engineer (Agentic AI) Intern

Gurugram, India

- Took LangGraph agents from local scripts to network-reachable production services holding long-term memory in Postgres, validating the self-hosted deployment architecture before the team committed to it.
- Benchmarked four LLM evaluation frameworks, DeepEval, Ragas, Langfuse and Arize Phoenix, and instrumented tracing and evaluation through LangSmith, so model regressions were caught in CI rather than reported by users.
- Root-caused a silent startup failure where an async Postgres driver deadlocked against Windows network I/O scheduling, and selected the foundation the rest of the migration was built on.

PROJECTS

CacheLLM, Semantic Caching Proxy for LLM APIs (Published on PyPI) | [Live](#) | [GitHub](#)

Python, FastAPI, Embeddings, Redis

- Built a drop-in semantic cache keeping the OpenAI request and response shape including streaming, routing each model to the host that serves it across OpenAI, Gemini, Groq, Claude, Grok, Ollama and Bedrock.
- Matched exact-hash first then cosine similarity over normalised embeddings, with the safe threshold calibrated per embedding model rather than assumed, reaching 77% hit rate and 78% cost reduction.
- Held zero wrong answers on new questions by namespacing on model, system prompt and sampling parameters, and refusing to cache any reply truncated by max_tokens.

Vayu, National Air-Quality Intelligence and Enforcement Platform (Deployed on Google Cloud Run) |

LightGBM, CNN-LSTM, DuckDB,

[Live](#) | [GitHub](#)

Python

- Enforcement targets the dirtiest sensor, not the site where action pays most. Forecast air quality with LightGBM on station data plus a CNN-LSTM over a 15,360-cell satellite grid, then ranked interventions by Gaussian-plume modelled return.
- Kept PyTorch out of the API request path entirely, training the CNN-LSTM offline and serving from DuckDB as one embedded columnar file, so scoring is fast without a model server in the loop.
- Verified that dispatched interventions actually worked using difference-in-differences against control regions, rather than reporting any post-intervention drop as proof of effect.

VaidyaMitra, Privacy-First Clinical GenAI Assistant | [Live](#) | [GitHub](#)

AWS Bedrock Nova, FastAPI, Next.js

- Masked every identifier before it reaches the model across a 38-endpoint API, so a hosted model never sees patient-identifying text.
- Benchmarked four Bedrock models on latency and multilingual quality to select Nova Pro at 605 ms against alternatives above 1,050 ms, and set maxTokens explicitly to avoid quota-reservation throttling.
- Matched prescriptions to Jan Aushadhi generics with substitution-safety warnings, surfacing savings up to 93.7% across 15 Indian languages.

TECHNICAL SKILLS

Generative AI: LangGraph, LangChain, LangSmith, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Agentic AI: LangGraph, MCP (Model Context Protocol), Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, LLMOps, DeepEval, Ragas, Langfuse, Arize Phoenix

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, Optuna, Pandas, NumPy, Feature Engineering, Cross-Validation, Calibration, SHAP, A/B Testing, Statistics

Backend and Infrastructure: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++, Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- Published three packages to PyPI, cachellm-proxy, creditsetu and medguardx-core, each installable, documented and covered by tests.
- KSP Datathon 2026 finalist: reached Level 2 with Kadi, a crime analytics platform built end to end over 50 days and deployed on Zoho Catalyst for the Karnataka State Police.

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 - 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (Apr 2025 - Mar 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.