

Adarsh Dwivedi

AI Product Manager

+91-9305597756 | 23ucs509@lnmiit.ac.in | [linkedin.com/in/adarshdwivedi30](https://www.linkedin.com/in/adarshdwivedi30) | github.com/adarshcod30

Final-year CS engineer who ships AI products with an abstention policy, an evaluation plan and an owner for every failure mode, not just a model. Built 32 products across 10+ national hackathons and datathons for users who cannot debug a model: police analysts, farmers and regulators. At Deloitte, recommended a self-hosted agent server over the commercial platform to keep agent state inside the firm, then proved it live on SAP at 8 of 8 tasks. Treats guardrails as product decisions.

EXPERIENCE

Deloitte India, Digital Excellence Centre <i>AI Engineer (Agentic AI) Intern</i>	Jun 2026 - Aug 2026 Gurugram, India
Evaluated three deployment routes against Deloitte's constraints and recommended Aegra: the commercial platform would move agent state outside the firm and an in-house runtime would rebuild a mature tool, while Aegra is open source on FastAPI and PostgreSQL.	
Proved the new path on the live SAP tenant: a Bank Accounting goal was planned into 8 configuration tasks, executed in a real browser through Playwright MCP, recovered from a duplicate-entry conflict and a mis-filled field, and finished 8 of 8 in about 15 minutes.	
Scoped the migration so every consumer kept working: the public API, business logic and front end stayed unchanged, and a single setting switches back to the original execution path with no redeployment.	

PROJECTS

Kadi, Crime Analytics Platform for the Karnataka State Police (<i>KSP Datathon 2026</i>) Live GitHub	Zoho Catalyst, Graph, Forecasting
Built for users who cannot debug a model: a station officer gets a ranked case queue with a plain-language reason, while the bureau reads state-wide analysis off the same graph of 59,985 FIRs.	
Set the responsible-AI boundary explicitly, ranking shared modus operandi as a hypothesis with a glass-box factor breakdown rather than naming a suspect.	
Shipped the full loop on Zoho Catalyst in 50 days, including an 0.870 AUC forecast, DBSCAN hotspots and evidence OCR, reaching Level 2 of the KSP Datathon.	
Vayu, National Air-Quality Intelligence and Enforcement Platform (<i>Deployed on Google Cloud Run</i>) Live GitHub	LightGBM, CNN-LSTM, Gemini
Designed an AI product whose output is an enforceable order, not a dashboard: rank interventions by modelled return, dispatch, then measure the result against control regions.	
Made an untrusted crowd signal usable by gating every citizen report through a vision model before it enters the record.	
KrishiMitra, Crop Intelligence Platform for Indian Smallholders (<i>Deployed, 12 Indian languages</i>) Live GitHub	CatBoost, CNN, FastAPI
Built for users who will not tolerate a confident wrong answer, framing 13 modules as questions a farmer actually asks and shipping in 12 Indian languages, free to use.	
Put a hard guardrail in front of the model, refusing any crop the district has no five-year record of growing, because a wrong sowing recommendation costs a season.	
AGENTIQ, Agentic API Testing and Security Validation (<i>B.Tech final-year project</i>) Live GitHub	MCP, LLM Agents, OpenAPI 3.1
Autonomy only sells if it is accountable. Framed an agentic testing platform around one rule: an agent may act, but only through 19 schema-validated tools, each granted per host and audited.	
Scoped it to the gap teams actually have, functional tests in one product and security tests in another, generating both from a URL and plain-English intent.	

TECHNICAL SKILLS

AI Product: AI Product Strategy, Evaluation and Acceptance Criteria, Model Selection and Cost Trade-offs, Human-in-the-Loop Design, Guardrails and Abstention Policy, Responsible AI and Governance, Prompt and Context Engineering, Retrieval (RAG) Design, Data Sourcing and Labelling, Latency and Unit Economics
Product: Product Strategy, Product Discovery, Problem Framing, PRD and Technical Specification, Roadmapping, Prioritisation (RICE), User Research, Metrics and KPI Definition, Adoption and Engagement, A/B Testing, Stakeholder Management, Competitive Analysis, Agile and Scrum, Go-to-Market
AI and ML Fluency: LangGraph, LangChain, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama
Data and Analytics: SQL, Pandas, NumPy, DuckDB, Statistical Testing, Experiment Design, H3 Geospatial Indexing, Graph Analytics, Power BI, Excel, Data Visualisation
Technical Fluency: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Redis, Zod Schema Validation

ACHIEVEMENTS

Competed in 10+ national hackathons and datathons: KSP Datathon 2026 finalist, Level 2 with Kadi for the Karnataka State Police; AWS AI for Bharat finalist with VaidyaMitra; Hackspllosion 2026 with Deloitte India, clearing Levels 1, 2 and 3; and SWITCH Energy-X, 11th of 93.
HackerRank Orchestrate: ranked #12 with a silver medal, top 1% (Sep 2026), up from #290 with a bronze medal, top 15% (Aug 2026), in 24-hour agentic AI challenges.
micro1 Frontier Engineering Challenge 2026: built a research-artifact verifier measured across 742 repositories.
32 public repositories, most deployed and free to use, with three packages published to PyPI and every reported number reproducible from the repository.

EDUCATION

The LNM Institute of Information Technology, Jaipur <i>B.Tech, Computer Science and Engineering</i>	2023 - 2027 CGPA: 7.11 / 10
---	---------------------------------------

LEADERSHIP

Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.