Adarsh Dwivedi

Data Analyst

+91-9305597756 | 23ucs509@lnmiit.ac.in | linkedin.com/in/adarshdwivedi30 | github.com/adarshcod30

Final-year CS engineer who turns fragmented data into something a decision-maker can act on. Joined 59,985 synthetic police records across 31 districts into one analysable register, reconciled 99.6% of a national aviation dataset without manual mapping and held forecast error to 11.1% median, and ranked 100,000 candidate profiles in under 18 seconds on a single CPU core. Every figure traces back to the source row it came from.

EXPERIENCE

Deloitte India, Digital Excellence Centre **Jun 2026 - Aug 2026**
AI Engineer (Agentic AI) Intern Gurugram, India

- Migrated SAP Agentic, Deloitte's system that automates SAP configuration, from LangGraph running inside a FastAPI request handler onto a self-hosted Aegra agent server, with threads, runs and checkpoints in PostgreSQL inside the firewall at no licence cost.
- Set the migration's acceptance test as field-by-field identical output between the original and migrated paths on the same input, rather than output that merely looked plausible, and held the public API unchanged so no consumer needed modification.
- Proved the new path on the live SAP tenant: a Bank Accounting goal was planned into 8 configuration tasks, executed in a real browser through Playwright MCP, recovered from a duplicate-entry conflict and a mis-filled field, and finished 8 of 8 in about 15 minutes.

Suganda Skincare, Growth Analysis (Contract) **Sep 2026 - Oct 2026**
Data and Growth Analyst, Freelance Remote

- Consolidated 19 Shopify exports into a DuckDB warehouse of 174,747 orders and 86,146 customers, fixing traps that each gave a wrong answer first: order money sat only on the first line row (AOV read Rs 587) and a payment name containing 'cash' misread 47,375 prepaid orders as COD.
- Beat a free benchmark on every model: all 10 revenue and demand series beat seasonal naive on rolling-origin backtests, return prediction reached 6.07x lift against 5.32x for an RFM score, and survival analysis corrected the one-year repurchase rate from a naive 32% to 57%.
- Sized causes against a counterfactual: customer count explains 114% of the rise and 89% of the fall in revenue, returns climb from 12.4% to 20.7% as despatch slows, and Tier 2 and smaller cities carry 53.8% of revenue.

PROJECTS

Kadi, Crime Analytics Platform for the Karnataka State Police *(KSP Datathon 2026)* | **Live** | **GitHub** Python, SQL, Entity Resolution, DBSCAN

- Joined 59,985 synthetic FIRs, generated on the real KSP schema, across 31 districts and 298 stations into one analysable register, with entity resolution collapsing duplicate offenders and name variants across station boundaries.
- Delivered per-capita ranking, rank-shift analysis, crime-mix by urbanisation band, DBSCAN hotspot detection and a state forecast at 0.870 AUC three months ahead.
- Surfaced 26,168 investigation-health exceptions as a ranked worklist with plain-language reasons, turning a reporting exercise into a queue a station can actually work through.

Air Cargo Intelligence, Agentic Analytics over Indian Aviation Data | **Live** | **GitHub** Python, PostgreSQL, SARIMA, Multi-Agent

- Indian air-cargo data is spread across DGCA and AAI releases in incompatible shapes, so built an ingestion layer into a PostgreSQL star schema with an idempotent loader, reconciling 99.6% of rows (12,740 of 12,790) without manual mapping.
- Competed five forecast candidates per series, naive, drift, recent-mean, seasonal-naive and SARIMA, picking the winner by rolling-origin backtest, reaching 11.1% median MAPE with 82.4% coverage on the 80% prediction interval.
- Refused documents rather than guessing, recording a reason for each of the 15 refused out of 261, and held anomaly alerts to 3.1 a month above a materiality bar so the output stays actionable.

Talent Intelligence Engine, Large-Scale Candidate Ranking and Verification | **GitHub** Python, Pandas, Streamlit

- Established first that per-profile LLM scoring was operationally infeasible at 100,000 candidates on cost and latency, then built the ranking to run in under 18 seconds on a single CPU core with zero network calls.
- Scored profiles across four weighted dimensions, role fit and skills at 30% each, experience and logistics at 20%, behind a deterministic filter that removed 100% of honeypot records before ranking ran.

TECHNICAL SKILLS

Data and Analytics: SQL, Pandas, NumPy, DuckDB, Statistical Testing, Experiment Design, H3 Geospatial Indexing, Graph Analytics, Power BI, Excel, Data Visualisation

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, Optuna, Pandas, NumPy, Feature Engineering, Cross-Validation, Calibration, SHAP, A/B Testing, Statistics

Backend and Tooling: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++, Java, JavaScript, TypeScript, SQL

Core Computer Science: Data Structures and Algorithms, OOP, DBMS, Operating Systems, Computer Networks, System Design

ACHIEVEMENTS

- KSP Datathon 2026 finalist: reached Level 2 with Kadi, a crime analytics platform built end to end over 50 days and deployed on Zoho Catalyst for the Karnataka State Police.
- 39 public repositories, most deployed and free to use, with every reported number reproducible from the repository and failures published beside results.

EDUCATION

The LNM Institute of Information Technology, Jaipur **2023 - 2027**
B.Tech, Computer Science and Engineering CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.