

Adarsh Dwivedi

Deep Learning Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | linkedin.com/in/adarshdwivedi30 | github.com/adarshcod30 | huggingface.co/adarshcod30

Final-year CS engineer training vision models that state their confidence and abstain when they should. Built a three-backbone deepfake ensemble with calibrated confidence and per-backbone Grad-CAM, and a retinopathy grader using rank-consistent ordinal heads on patient-level splits, with localisation measured against annotated lesions rather than asserted. Models are deployed to Hugging Face and Streamlit, with scope limits published beside the results.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 – Aug 2026

AI Engineer (Agentic AI) Intern

Gurugram, India

- Took LangGraph agents from local scripts to network-reachable production services holding long-term memory in Postgres, validating the self-hosted deployment architecture before the team committed to it.
- Benchmarked four LLM evaluation frameworks, DeepEval, Ragas, Langfuse and Arize Phoenix, and instrumented tracing and evaluation through LangSmith, so model regressions were caught in CI rather than reported by users.
- Root-caused a silent startup failure where an async Postgres driver deadlocked against Windows network I/O scheduling, and selected the foundation the rest of the migration was built on.

PROJECTS

OpenForensics, Deepfake Face Detection Ensemble | [Live](#) | [GitHub](#)

PyTorch, ResNet50, VGG16, EfficientNetV2, Grad-CAM

- A detector that outputs only a label is unusable as evidence, so built a three-backbone ensemble of ResNet50, VGG16 and EfficientNetV2-B0 with calibrated confidence and an explicit borderline warning band.
- Trained on a two-stage schedule and served per-backbone Grad-CAM, so a reviewer sees which region drove each model's vote instead of trusting the aggregate score.
- Published the training, evaluation and robustness evidence behind every number in an eight-page dashboard, and stated the scope limits plainly: face crops only, swaps rather than generated images, and a single dataset.

Diabetic Retinopathy Screening, Explainable Grading for Rural Screening Programmes | [Live](#) | [GitHub](#)

PyTorch, Ordinal Heads, Grad-CAM, CLAHE

- Grading a 0 as a 4 is not the same mistake as grading a 0 as a 1, so trained ICDR grades 0 to 4 with rank-consistent ordinal heads instead of plain cross-entropy, on patient-level splits where both eyes never straddle train and test.
- Gated every image through a 3-class gradability model fused with focus, illumination and field-of-view features, returning why an image was rejected so the operator can recapture rather than guess.
- Detected sub-pixel microaneurysms with morphological top-hat and matched filtering before CNN classification, and measured Grad-CAM localisation against annotated lesions rather than asserting the heatmaps meant anything.

Vayu, National Air-Quality Intelligence and Enforcement Platform (Deployed on Google Cloud Run) | [Live](#) | [GitHub](#)

PyTorch, CNN-LSTM, Satellite Imagery

- Trained a CNN-LSTM over a 15,360-cell satellite grid for spatio-temporal air-quality forecasting, fused with LightGBM on ground-station readings so neither source has to be trusted alone.
- Kept the trained network out of the request path, training offline and serving precomputed scores from an embedded DuckDB file, removing model-server latency from every user request.

TECHNICAL SKILLS

Deep Learning: PyTorch, TensorFlow/Keras, CNNs, ResNet50, VGG16, EfficientNetV2, CNN-LSTM, Transformers, Transfer Learning, Ordinal Regression, Grad-CAM, OpenCV

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, Optuna, Pandas, NumPy, Feature Engineering, Cross-Validation, Calibration, SHAP, A/B Testing, Statistics

Generative AI: LangGraph, LangChain, LangSmith, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Backend and Deployment: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++ , Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- SWITCH Energy-X (India): 11th of 93 on the private leaderboard, recovering a hidden physical law from 500,000 unlabeled, 26%-missing sensor rows.
- KSP Datathon 2026: reached Level 2 with Kadi, a crime analytics platform built end to end over 50 days and deployed on Zoho Catalyst for the Karnataka State Police.

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 – 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (Apr 2025 – Mar 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 – Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.