

Adarsh Dwivedi

DevOps and Platform Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | [linkedin.com/in/adarshdwivedi30](https://www.linkedin.com/in/adarshdwivedi30) | github.com/adarshcod30 | pypi.org/user/adarshcod30

Final-year CS engineer who treats reproducibility as the deliverable. Built Dockerized pipelines where the same inputs always produce the same outputs, a caching proxy on PyPI that fails open and reports itself degraded rather than failing requests, and services hardened with JWT RBAC that refuse to start without their encryption key. Deploys across AWS, Cloud Run, Vercel and Render, with CI that fails loudly rather than silently.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 - Aug 2026

AI Engineer (Agentic AI) Intern

Gurugram, India

- Replaced request-time graph assembly with a graph factory that registers one static graph and carries per-request variation in run state, verifying the migrated output identical to the original field by field behind an unchanged API.
- Held the public API fixed so the front end and test collections needed no change, and made the migration reversible through one setting, EXECUTION_BACKEND, keeping the original execution path byte for byte so rollback needs no redeploy.
- Docker was unavailable, so worked out and documented a Docker-free deployment of FastAPI, Aegra and Playwright MCP on a shared Linux server against PostgreSQL, after tracing unrelated-looking failures to an unsupported runtime and rebuilding on Python 3.13.

PROJECTS

PrimeTradeML, Reproducible MLOps Batch Pipeline | [GitHub](#)

Docker, Python, CI, YAML

- Built a deterministic, Dockerized batch pipeline for trading-signal generation where the same inputs always reproduce the same outputs, because a signal you cannot reproduce is a signal you cannot audit.
- Wired validation and observability into CI so a run fails loudly on a data or schema problem rather than emitting a quietly wrong signal, with per-stage timings recorded on every run.
- Kept configuration in YAML and the environment in the image, so a run is defined entirely by its inputs and its tag.

CacheLLM, Semantic Caching Proxy for LLM APIs (Published on PyPI) | [Live](#) | [GitHub](#)

Python, FastAPI, Docker, Redis, PyPI

- Shipped a caching proxy to PyPI fronting OpenAI, Gemini, Groq, Claude, Grok, Ollama and Bedrock behind one interface, routing each model to the host that serves it from that host's live model list.
- Made it run with nothing installed: the cache is an in-process numpy matrix by default with Redis optional, and the proxy fails open, forwarding upstream and reporting itself degraded rather than failing the request.
- Measured a 77% hit rate and 78% cost reduction, with a per-request policy deciding whether a response is safe to cache at all and for how long.

MedGuardX, Context-Aware PII and PHI Masking (Published on PyPI) | [Live](#) | [GitHub](#)

Python, FastAPI, Docker, JWT RBAC

- Shipped three layers, a stateless engine on PyPI, a hardened FastAPI service and a Next.js app, with the spaCy model size configurable so an operator picks the accuracy and memory trade-off at deploy time.
- Enforced authentication on every protected route with the role read from the JWT rather than the body, and refused privileged self-registration so admins are seeded out of band.
- Required the encryption key in production and made the service fail fast when it is missing, rather than silently falling back to a throwaway key.

TECHNICAL SKILLS

Infrastructure and DevOps: Docker, AWS, Google Cloud Run, Vercel, Render, GitHub Actions, CI/CD, Linux, Git, Observability, Reproducible Builds

Backend and APIs: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++, Java, JavaScript, TypeScript, SQL

Core Computer Science: Data Structures and Algorithms, OOP, DBMS, Operating Systems, Computer Networks, System Design

ACHIEVEMENTS

- Published three packages to PyPI, cachellm-proxy, creditsetu and medguardx-core, each installable, documented and covered by tests.
- 32 public repositories, most deployed and free to use, with every reported number reproducible from the repository and failures published beside results.

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 - 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.