Adarsh Dwivedi

Generative AI Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | linkedin.com/in/adarshdwivedi30 | github.com/adarshcod30

Final-year CS engineer who builds generative systems that refuse rather than guess. At Deloitte, migrated a LangGraph agent onto a self-hosted Aegra server behind an unchanged FastAPI interface and proved it live on SAP, completing 8 of 8 tasks. Builds retrieval-first: identifiers masked before the model sees them, a vision-language verifier hardened against prompt injection inside images, and an agent where no product answer may originate from the model.

EXPERIENCE

Deloitte India, Digital Excellence Centre **Jun 2026 - Aug 2026**

Full Stack AI Engineer Intern Gurugram, India

- Migrated SAP Agentic, Deloitte's system that automates SAP configuration, from LangGraph running inside a FastAPI request handler onto a self-hosted Aegra agent server, with threads, runs and checkpoints in PostgreSQL inside the firewall at no licence cost.
- Proved the new path on the live SAP tenant: a Bank Accounting goal was planned into 8 configuration tasks, executed in a real browser through Playwright MCP, recovered from a duplicate-entry conflict and a mis-filled field, and finished 8 of 8 in about 15 minutes.
- Built capability in five layers on a multi-agent travel planner, from a single agent to a supervisor with three specialists, a human approval gate, real browser tools over MCP and serving through Aegra, proving each layer before starting the next.

PROJECTS

VaidyaMitra, Privacy-First Clinical GenAI Assistant | **Live** | **GitHub** AWS Bedrock Nova, FastAPI, Next.js

- Clinical text cannot be handed to a hosted model as-is, so every identifier is masked before the request leaves the service, enforced across a 38-endpoint API rather than at one choke point.
- Benchmarked nine Amazon Bedrock models on latency and multilingual quality and selected Nova Pro at 605 ms, setting maxTokens on every call to avoid quota-reservation throttling.

Multi-Modal Evidence Review, Vision-Language Claim Verification *(Orchestrate Jun 2026)* | **GitHub** Gemini Vision, Python

- A blurry photo of the right bumper, a sharp photo of the wrong bumper and a sharp photo of real damage look identical to a naive classifier. Built a four-stage vision-language pipeline that separates them.
- Analysed each image in its own call rather than batching, keeping every visual judgement independent, and defended against prompt injection arriving as rendered text inside the submitted photos.
- Kept the risk engine pure Python with zero model calls so aggregation cannot be talked out of a verdict, reached 75.0% on the headline decision, and published that n=20 makes every figure low-resolution.

AutoStream, RAG-Grounded Conversational Lead Agent *(Deployed)* | **Live** | **GitHub** LangGraph, Gemini, RAG

- A sales agent that invents a product detail costs a customer, so no product answer comes from the model directly: every one is routed through retrieval, and greetings bypass retrieval entirely.
- Built the conversation as a LangGraph state machine, intent to retrieval to lead capture to response, with a typed state carrying intent, stage and extracted fields across turns.

Message Notification Router, Message Triage Agent *(HackerRank Orchestrate Aug 2026, ranked #290)* | **GitHub** Python, Claude, Gemini

- Let the evaluation loop find the bug: both models scored 86.7% and failed the same way, which pointed at the prompt and not the model. Changing only the briefing lifted action accuracy to 93.3% and cut Brier from 0.119 to 0.071 on the same rows, for $0.033.
- Stayed correct through a real quota outage: with the API returning 429 on every request, a circuit breaker produced all 110 decisions in 19.8 s, matching the full run 110 of 110, and the write-up states that n = 30 makes one row worth 3.3 points.

TECHNICAL SKILLS

Generative AI: LangGraph, LangChain, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Agentic AI: LangGraph, MCP (Model Context Protocol), Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, Agent Protocol, Aegra, Playwright MCP, LLMOps

LLM Reliability: Structured Output, Schema Validation, Guardrails, Fallbacks and Circuit Breakers, Evaluation Harnesses, Calibration, Prompt-Injection Defence, AI-Assisted Development with transcript logging

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, Optuna, Pandas, NumPy, Feature Engineering, Cross-Validation, Calibration, SHAP, A/B Testing, Statistics

Backend and Infrastructure: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++, Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- HackerRank Orchestrate: ranked #12 with a silver medal, top 1% (Sep 2026), up from #290 with a bronze medal, top 15% (Aug 2026), in 24-hour agentic AI challenges.
- Competition submissions, each public on GitHub: Buy or Wait, #12 · Message Notification Router, #290 bronze · Multi-Modal Evidence Review, #427 · Artifact Repro Triage, micro1

EDUCATION

The LNM Institute of Information Technology, Jaipur **2023 - 2027**

B.Tech, Computer Science and Engineering CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.