

Adarsh Dwivedi

Generative AI Engineer

+91-9305597756 | 23ucs509@lnmiit.ac.in | [linkedin.com/in/adarshdwivedi30](https://www.linkedin.com/in/adarshdwivedi30) | github.com/adarshcod30

Final-year CS engineer who builds generative systems that refuse rather than guess. At Deloitte, worked full stack across FastAPI services and LLM tool-calling pipelines, and benchmarked four evaluation frameworks so regressions surfaced in CI rather than from users. Builds retrieval-first: identifiers masked before the model sees them, a vision-language verifier hardened against prompt injection inside images, and an agent where no product answer may originate from the model.

EXPERIENCE

Deloitte India, Digital Excellence Centre

Jun 2026 - Aug 2026

Full Stack AI Engineer Intern

Gurugram, India

- Took LangGraph agents from local scripts to network-reachable production services holding long-term memory in Postgres, validating the self-hosted deployment architecture before the team committed to it.
- Benchmarked four LLM evaluation frameworks, DeepEval, Ragas, Langfuse and Arize Phoenix, and instrumented tracing and evaluation through LangSmith, so model regressions were caught in CI rather than reported by users.
- Root-caused a silent startup failure where an async Postgres driver deadlocked against Windows network I/O scheduling, and selected the foundation the rest of the migration was built on.

PROJECTS

VaidyaMitra, Privacy-First Clinical GenAI Assistant | [Live](#) | [GitHub](#)

AWS Bedrock Nova, FastAPI, Next.js

- Clinical text cannot be handed to a hosted model as-is, so every identifier is masked before the request leaves the service, enforced across a 38-endpoint API rather than at one choke point.
- Benchmarked four Amazon Bedrock models on latency and multilingual quality and selected Nova Pro at 605 ms against alternatives above 1,050 ms, setting maxTokens on every call to avoid quota-reservation throttling.
- Matched prescriptions to Jan Aushadhi generics with substitution-safety warnings, surfacing savings up to 93.7%, and simplified clinical reports into 15 Indian languages.

Multi-Modal Evidence Review, Vision-Language Claim Verification (Orchestrate Jun 2026) | [GitHub](#)

Gemini Vision, Python

- A blurry photo of the right bumper, a sharp photo of the wrong bumper and a sharp photo of real damage look identical to a naive classifier. Built a four-stage vision-language pipeline that separates them.
- Analysed each image in its own call rather than batching, keeping every visual judgement independent, and defended against prompt injection arriving as rendered text inside the submitted photos.
- Kept the risk engine pure Python with zero model calls so aggregation cannot be talked out of a verdict, reached 75.0% on the headline decision, and published that n=20 makes every figure low-resolution.

AutoStream, RAG-Grounded Conversational Lead Agent (Deployed) | [Live](#) | [GitHub](#)

LangGraph, Gemini, RAG

- A sales agent that invents a product detail costs a customer, so no product answer comes from the model directly: every one is routed through retrieval, and greetings bypass retrieval entirely.
- Built the conversation as a LangGraph state machine, intent to retrieval to lead capture to response, with a typed state carrying intent, stage and extracted fields across turns.
- Classified intent with an LLM and rule-based checks together rather than prompting alone, so a misread message fails to a safe branch, and covered the graph with 55 tests before deploying it live.

TECHNICAL SKILLS

Generative AI: LangGraph, LangChain, LangSmith, RAG, Prompt Engineering, Vector Search, Embeddings, Evaluation and Calibration, Amazon Bedrock, Gemini, OpenAI, Groq, Ollama

Agentic AI: LangGraph, MCP (Model Context Protocol), Multi-Agent Orchestration, Tool Calling, Human-in-the-Loop, LLMops, DeepEval, Ragas, Langfuse, Arize Phoenix

Machine Learning: Scikit-learn, XGBoost, LightGBM, CatBoost, Optuna, Pandas, NumPy, Feature Engineering, Cross-Validation, Calibration, SHAP, A/B Testing, Statistics

Backend and Infrastructure: FastAPI, Pydantic, Node.js, Express, REST APIs, PostgreSQL, MongoDB, DuckDB, Next.js, React, Streamlit

Languages: Python, C/C++ , Java, JavaScript, TypeScript, SQL

ACHIEVEMENTS

- HackerRank Orchestrate: ranked #12 with a silver medal, top 1% (Sep 2026), up from #290 with a bronze medal, top 15% (Aug 2026), in 24-hour agentic AI challenges.
- Competition submissions, each public on GitHub: [Buy or Wait, #12](#) · [Message Notification Router, #290 bronze](#) · [Multi-Modal Evidence Review, #427](#) · [Artifact Repro Triage, micro1](#)

EDUCATION

The LNM Institute of Information Technology, Jaipur

2023 - 2027

B.Tech, Computer Science and Engineering

CGPA: 7.11 / 10

LEADERSHIP

- Coordinator, Debate Society, LNMIIT Jaipur (May 2025 - Apr 2026): led a 150+ member society across debates, discussions and inter-college competitions.
- Secretariat and Organizer, LNMIIT Model United Nations (Aug 2025 - Jan 2026): organised a large-scale MUN conference, coordinating delegates, committees, schedules and logistics.